ENHANCEMENT TO XOS OPERATING SYSTEM AND XFS FILE SYSTEM

A THESIS

Submitted by

KRUTHIKA SURESH VED B110300CS
SIKHA V MANOJ B110572CS *and*
SONIA V MATHEW B110495CS

In partial fulfilment for the award of the degree of

BACHELOR OF TECHNOLOGY
IN
COMPUTER SCIENCE AND ENGINEERING

Under the guidance of
DR. K MURALIKRISHNAN

DEPARTMENT OF COMPUTER ENGINEERING
NATIONAL INSTITUTE OF TECHNOLOGY CALICUT
NIT CAMPUS PO, CALICUT
KERALA, INDIA 673601

May 14, 2015

ACKNOWLEDGEMENTS

DECLARATION

"I hereby declare that this submission is my own work and that, to the best of my knowledge and belief, it contains no material previously published or written by another person nor material which has been accepted for the award of any other degree or diploma of the university or other institute of higher learning, except where due acknowledgment has been made in the text".

Place: Signature :
Date: Name :
                                          Reg.No:

                                          Signature :
                                          Name :
                                          Reg.No:

                                          Signature :
                                          Name :
                                          Reg.No:

CERTIFICATE

This is to certify that the thesis entitled: "ENHANCEMENT TO XOS OP-ERATING SYSTEM AND XFS FILE SYSTEM" *submitted by Sri/Smt/Ms* KRUTHIKA SURESH VED B110300CS , SIKHA V MANOJ B110572CS, *and* SONIA V MATHEW B110495CS *to National Institute of Technology Calicut towards partial fulfillment of the requirements for the award of Degree of Bachelor of Technology in Computer Science Engineering is a bonafide record of the work carried out by him/her under my/our supervision and guidance.*

Signed by Thesis Supervisor(s) with name(s) and date

Place:
Date:

Signature of Head of the Department

Office Seal

Contents

Chapter

Abstract

eXpOS or eXperimental Operating System is a tiny multiprogramming operating system. It has a very simple specification that allows a junior undergraduate computer science student to implement it in a few months, subject to availability of adequate hardware and programming platform support. It is also an instructional tool to learn and implement OS data structures and functionalities on a simulated machine called XSM (eXperimental String Machine).

Figures

Figure

Chapter 1

Problem Definition

This project aims to modify and update the XOS operating sytem by

- Re-engineering the eXpFS system.

- Including inter process communication

- Redesigning process model

- Introducing shared memory model

- Adding more system calls thus broadening the features

The updated version is named as eXpOS/Experimental Operating System. Thus, adding the above features, eXpOS would become more advanced and closer to operating systems that are available in the market.

Chapter 2

Literature Survey

There are several instructional operating systems available like Nachos, Topsy etc. developed by various universities.

Not Another Completely Heuristic Operating System, or Nachos, is one such instructional software developed at the University of California, Berkeley. Nachos supports threads, user-level processes, virtual memory and interrupt-driven input output devices. The only difference between Nachos and a real operating system is that Nachos runs as a single Unix process. It also supports synchronization and networking.

Low-level Nachos routines frequently disable and re-enable interrupts to achieve mutual exclusion. Synchronization facilities are provided through semaphores. Nachos assumes that only a single user program exists at a given time. There are two versions of the Nachos filesystem. A stub version is simply a front-end to the Unix filesystem, so that users can access files within Nachos without having to write their own file system. The second version allows students to implement a filesystem on top of a raw disk, which is in fact a regular Unix file that can only be accessed through a simulated disk. The simulated disk accesses the underlying Unix file exclusively through operations that read and write individual sectors. Both file systems provide the same service and interface.

Similar to Nachos, Topsy, or Teachable Operating System, is an instructional tool developed by ETH-Zurich. In Topsy all threads are running in one address space and may share global memory between them. Synchronization of shared memory is accomplished via messages (AOSC, section 4.5.1AOSC) between multiple threads. The private area of a thread is its stack which is not protected against (faulty) accesses from other threads. Topsy divides the memory into two address spaces: one for user threads and the other for the OS kernel (separation of user and kernel process). , the two address spaces are embedded in one virtual address space, although no swapping of pages to secondary memory is supported. The memory is organized in paged manner. It provides a framework for writing and installing hardware device drivers. The kernel contains three main modules reflecting the basic OS tasks: the memory manager, the thread manager and the I/O subsystem.

Almost every other instructional operating system provides a skeleton of an operating system. However in eXpOS, only the specification has been laid out, and students learn to implement XOS from ground up using the tools provided. In this project, a simple high level language called ExpL (Experimental Language) and its cross-compiler to XSM instruction set is provided to write user programs to test eXpOS. An XSM dependent language called SPL (System Programmers Language) and its cross-compiler is provided to program the OS itself. Most instructional operating systems use the UNIX filesystem for file management by the operating system. Instead, eXpOS provides a native file system known as eXpFS (Experimental File System).

Chapter 3

Introduction

Project eXpOS or experimental Operating System is a educational platform to develop an operating system. Several instructional operating systems like Nachos, Pintos, GeekOS etc. have been developed by various universities for this purpose. Nachos has been one of the most popular instructional operating systems available and is being used in many institutes across the world. XOS, the previous version of eXpOS, was such a tool that helps undergraduate computer science students acquire an elementary understanding of the practical aspects of an operating system.

XOS had features like multiprogramming, process management, a primitive filesystem and virtual memory. In order to incorporate simplicity, OS features like inter-process communication, device management, file caching, file permissions etc. had been excluded from XOS. In the present version, eXpOS, some of these features such as inter process communication, asynchronous disk access, heap memory, etc had been included to make the student more proficient in Operating System concepts. The basic Machine model consists of memory, disk and the CPU. eXpOS includes specification for a set of system calls, the scheduler, the exception handler, device drivers etc. Further hardware support required includes the timer, disk controller, Input-output system etc.

The primary components of the project include a simulated machine hardware (XSM), file system (eXpFS) and the operating system (eXpOS). No code base for the operating system is provided and the operating system is completely implemented by the student. Apart from the primary components, various tools are provided as part of the development environment. They include languages like Experimental Programming Language (ExpL) and System Programmer's Language (SPL) and their cross compilers to XSM instruction set, XSM debugger, and a UNIX-XFS interface to transfer files between a UNIX machine and the XFS disk (the eXpFS disk is itself a UNIX file).

Chapter 4

Experimental File System(eXpFS)

4.1   Introduction

The eXpFS logical file system provides a file abstraction that allows application programs to think of each data (or executable) file stored in the disk as a continuous stream of data (or machine instructions) without having to worry about the details of disk block allocation. eXpOS provides a sequence of file system calls through which application programs can create/read/write files. These system calls are OS routines that does the translation of the user request into physical disk block operations.

In addition to the eXpOS system call interface, the eXpFS specification also requires that there is an external interface through which executable and data files can be loaded into the file system externally. The external interface for eXpOS implementation on the XSM machine is the XFS Interface.

This section discusses the abstract logical view provided by eXpFS to the eXpOS application programmer.

4.2   eXpFS File system organization

The eXpFS logical file system comprises of files organized in a single directory called the root. The root is also treated conceptually as a file. Every eXpFS file is a sequence of words. Associated with each eXpFS file there are three attributes

- name, size and type, each attribute being one word long. The filename must be a string. Each file must have a unique name. The size of the file will be the total number of words stored in the file.(The maximum size of a file is operating system dependent). There are three types of eXpFS files - the root, data files and executable files. Each file in eXpFS has an entry in the root called its root entry.

4.2.1   The eXpFS Root File

The root file has the name root and contains meta-data about the files stored in the file system. For each file stored in eXpFS, the root stores three words of information - filename, file-size and file-type. This triple is called the root entry for the file. The first root entry is for the root itself. The order in which the remaining entries appear is not specified and can vary with the implementation. (The maximum size of the root file is defined by XFS_ROOTSIZE.).

Example: If the file system stores two files - a data file, file.dat, of size 700 words and an executable file, confirm this program.xexe, of 1025 words, the root file will contain the following information.

File name	File size	File Type
root	512	ROOT
file.dat	1024	DATA
program.xexe	1536	EXEC

Figure 4.1: Root File Entry

The operations on the root file are Open, Close, Read and Seek.

4.2.2   eXpFS Data files

A data file is a sequence of words. The maximum number of words permissible in a file is defined by the constant MAX_FILE_SIZE. (It is a recommended programming convention to use the extension ".dat" for data files). eXpFS treats

every file other than root and executable files as a data file. The Create system call automatically sets the file type field in the root entry for any file created through the create system call to DATA.

eXpOS allows an application program to perform the following operations on data files: Create, Delete, Open, Close, FLock, FUnlock, Read, Write, Seek. Application programs can create only data files using the Create system call. In addition to this, data files may be loaded into the eXpFS file system using the external interface.

4.2.3   eXpFS Executable files

These contain executable code for programs that can be loaded and run by the operating system. From the point of view of the eXpFS file system alone, executable files are just like data files except that file type is EXEC in the root entry. eXpFS specification does not allow executable files to be created by application programs. They can only be created externally and loaded using the external interface. However, application programs can read or modify executable files like data files.

Executable files are essentially program files that must be loaded and run by the operating system. Hence the Operating system imposes certain structure on these files (called the executable file format). Moreover, the instructions must execute on the machine on which the OS is running. Thus, there is dependency on the hardware as well. Typically, an application program written in a high level language (like ExpL) is compiled using a compiler that generates the executable file. The compiler generates executable file that is dependent on the operating system as well as the target machine.

An OS implementation on a particular machine specifies an application binary interface (ABI). The eXpOS ABI for XSM machine is described in Appendix A.

Application programs are typically written in a high level language like ExpL. The ExpL compiler for eXpOS running on the XSM machine generates target code based on the ABI specification for eXpOS on XSM.

The executable file format recognized by eXpOS is called the Experimental executable file (XEXE) format. In this format, an executable file is divided into two sections. The first section is called header and the second section called the code (or text) section. The code section contains the program instructions. The header section contains information like the size of the text and data segments in the file, the space to be allocated for stack and heap areas when the program is loaded for execution etc. This information is used by the OS loader to map the file into a virtual address space and create a process in memory for executing the program. An application program can read/modify an executable file like a data file using the standard system calls on data files. In addition to this, the Exec system call invokes the eXpOS loader that loads an executable file in XEXE format into memory for execution.

Chapter 5

Process Model

A program under execution is called a process. A process is newly created when a process already in execution invokes the Fork system call. The first process, the INIT process, is created by the OS during bootstrap by loading a code stored in a pre-defined disk location to memory and setting up a process. The OS assigns a unique integer identifier called process id for each process when it is created. The process id does not change during the lifetime of the process. The process that creates the new process is called the parent process of the newly created process.

5.1   The Structure of Processes

eXpOS associates a virtual (memory) address space for each process.

Figure 5.1: Process Model in eXpOS

The address space of a process is a contiguous sequence of memory ad-

dresses, starting from zero, accessible to a process. The eXpOS logically partitions the address space into four regions library, heap, code and stack. These regions are mapped into physical memory using hardware mechanisms like paging/segmentation.

Every process corresponds to some executable file stored in XEXE format stored in the eXpFS file system. The XEXE header of the executable file contains the information about how much space must be allocated for the various memory regions. When the program is loaded into memory by the operating system, the OS reads the header and sets up the regions of the virtual address space accordingly. Once the layout of the virtual address space is clear, the OS maps the virtual address space into the physical memory. The part of the OS which does all these tasks is called the OS loader. The eXpOS loader is the interrupt service routine corresponding to the Exec system call.

A process may open files or semaphores. The OS associates a file handle with each open instance of a file. Similarly, the OS assigns a semaphore identifier (semid) for each semaphore acquired by the process. The file handles and semids acquired by a process are also attributes of a process.

Note: In addition to the above attributes of a process that are visible to application/system programs, a process under execution at any given point of time has an execution context. The context of a process refers to the contents of the registers, instruction pointer, contents of the memory etc. These are hardware dependent and are managed internally by the OS. In fact, managing the execution contexts of multiple processes simultaneously and running them all in one machine is the main challenge in the design and implementation of a multiprogramming OS. However, the OS hides these internal details from the application programs as well as system programs like compilers.

5.2   Operations on Processes

The two most fundamental operations associated with process are Fork and Exec. The remaining operations are Exit, Wait, Signal, Getpid and Getppid.

5.2.1   Semantics of Exec operation

(1) The OS closes all files and semaphores opened by the process. A new address space is created replacing the existing one. The new process inherits the process id of the calling process.

(2) The code (and static data, if any) of the executable file are loaded into the code (and stack) regions of the new address space. The system library is mapped to the library region and stack is initialized to empty.

(3) The machine instruction pointer is set to the location specified in the executable header. The machine stack pointer is initialized to the beginning of the stack. From here, execution continues with the newly loaded program.

5.2.2   Semantics of Fork operation

(1) A new child process with a new process id and address space is created which is an exact replica of the original process with the same library, code,stack and heap regions. (The OS assigns a new process id for the child and returns this value to the parent as the return parameter of the fork system call.) The heap, code and library regions of the parent are shared by the child. Stack is separate for the child and is not shared.

(2) All open file handles and semaphores are shared by the parent and the child. This means that if one of the processes write/read into/from or adjust the file handle using the lseek system call, the corresponding file

handle of the other process also gets automatically updated. Note that file handles (or semaphore identifiers) of files (or semaphores) that are opened (or created) subsequent to the fork operation by the parent or the child will be exclusive to the particular process and will not be shared. The parent and the child continue execution from here on.

5.2.3 Other Operations

The Exit system call terminates a process after closing all files and semaphores. The Wait system call suspends the execution of a process till another process exits or executes a Signal system call. The Signal system call resumes the execution of a process that was suspended by wait.

A process can get its process id using the Getpid system call. The pid of the parent process can be obtained using the Getppid system call.

5.3 Special Processes in eXpOS

eXpOS specifies two special processes, the idle and the init process. These are stored in a predefined location in the disk and loaded to the memory by the bootstrap loader. The main purpose of idle process is to run as a background process in an infinte loop. This is demanded by the OS so that the scheduler will always have a process to schedule. The init process is the first process executed by the OS. The process identifiers for the idle and init processes are fixed as 0 and 1 respectively.

A shell is an ExpL program which takes the name of an executable file as input and executes it. The shell process Forks itself and the child process invokes the Exec system call with the executable file as argument. The shell runs until the user stops the process.

Chapter 6

Synchronization and Access Control

6.1   Introduction

eXpOS assumes a single processor multi programming environment. This means that all processes exist concurrently in the machine and the OS time shares the machine between various processes. The OS specification requires that Round Robin scheduling is used with co-operative time-sharing. The OS does not provide any promise to the application program about the order in which processes will be executed.

However, application programs often need to stop and wait for another process to execute certain operations before proceeding. The OS provides system calls that allow user processes to synchronize execution.

6.2   Process Synchronization

eXpOS provides the Wait and Signal system calls for process synchronization. When a process executes the wait system call specifying the process id of another process as argument, the OS puts the calling process to sleep. This means, the OS will schedule the process out and won't execute it until one of the following events occur:

- The process specified as argument terminates by the exit system call.

- The process specified as argument executes a Signal system call.

6.3 Access Control

A second major concurrency related requirement is that when multiple processes access the same data (in memory or files), it is often required to have some kind of locking mechanism to ensure that when one process is accessing the shared data, no other process is allowed to modify the same. This is to ensure data integrity and this issue is called the critical section problem.

eXpOS provides (binary) semaphores to allow application programs to handle the critical section problem. A (binary) semaphore is conceptually a binary-valued variable whose value can be set or reset only by the OS through designated system call. (Internally how it is implemented is an OS concern and is oblivious to the application programmer).

A process can acquire a semaphore using the Semget system call, which returns a unique semid for the semaphore. The semaphore is initially not set. A process may open several files and semaphores and execute the Fork system call multiple times to create an several process that shares all the opened file handles, the heap memory region and all the semaphores acquired by the parent.

Any of these processes sharing a semaphore identifier can set the semaphore (called locking) using the SemLock system call giving the semid as an argument. SemUnLock will reset (called release) the semaphore waking up all other processes which went to sleep trying to lock the semaphore.

Just like semaphores, files can also be locked using the FLock and FUnLock system calls. They are discussed here.

Semantics of Locking operation for files and semaphores:

- If the semaphore/file is already locked, the process goes to sleep and wakes up only when the semaphore/file is free.

- Otherwise, process locks the semaphore/file and continues execution.

Chapter 7

Hardware Interrupts and Exception Handlers

7.1   Introduction

When the machine is powered on, the system is configured to start executing a ROM code in privileged mode. This code is called the bootstrap loader. This ROM code loads the first block of the disk into a pre-defined area in memory and transfers control to the newly loaded code. This code is called the OS startup code.

The OS startup code loads the system call routines into memory. These routines are loaded as software interrupt handlers. In addition to these, there are three hardware interrupt / exception handler modules that are to be loaded - the timer interrupt handler, the exception handler, the disk interrupt handler and the terminal handler. If the architecture supports other devices, then the corresponding device interrupt handlers also must be loaded by the OS Startup code.

7.2   Timer interrupt handler

The hardware requirement specification for eXpOS assumes that the machine is equipped with a timer device that sends periodic hardware interrupts. The OS scheduler is invoked by the hardware timer interrupt handler. eXpOS specification suggest that a co-operative multitasking round robin scheduling is employed. This means that a round robin scheduling is employed, but a process may go to sleep

inside a system call when:

- The resource which the process is trying to access (like a file or semaphore) is locked by another process (or even internally locked by another OS system call in concurrent execution).

- There is a disk or I/O device access in a system call which is slow. If the wait for the device access is to be avoided, there must be hardware support from the device to send a hardware interrupt when device operation is finished. This allows the OS to put the process on sleep for now, continue scheduling the remaining processes in round robin fashion and then wake up the sleeping process when the device sends the interrupt.

7.3   Exception handler

If a process generates an illegal instruction, an invalid address (outside its virtual address space) or do a division by zero (or other faulty conditions which are machine dependent), the machine will generate an exception. The exception handler must terminate the process, wake up all processes waiting for it (or resources locked by it) and invoke the scheduler to continue round robin scheduling the remaining processes.

The exception handler is invoked when a page required by the process is not present in the memory. This condition is known as a page fault. The module which handles demand paging (if the machine hardware supports demand paging) is invoked by the exception handler when there is a page fault. eXpOS specification does not require implementation of demand paging. However, most machines (including XSM) are equipped with hardware support for demand paging and using the feature can improve machine throughput considerably.

7.4   Disk Controller

eXpOS treats the disk as a special block device and assumes that the hardware provides low level block transfer routines to transfer disk blocks to memory (pages) and back. The block transfer routines contain instructions to initiate block-memory transfer by the disk controller hardware. After initiating the disk-memory transfer, the block transfer routine normally returns to the calling program, which sleeps for the disk operation to complete.

When the disk-memory transfer is complete, the disk controller raises a hardware interrupt. The interrupt service routine (handler) must be part of the OS code to be set up during the bootstrap. The disk interrupt handler is responsible for waking processes that went into sleep awaiting completion of the disk operation.

7.5   Terminal and other Device Handlers

All other data handling devices (other than the disk) are treated as stream devices. This means that each device allows transfer of only one word from memory to the device or back at a time. Some devices may permit only write (like a printer) whereas some devices may permit only read. It is assumed that for each device there are associated low level routines that can be invoked by the OS to transfer data and control instructions. Some of these devices may raise a hardware interrupt when the transfer is complete. Thus, for each device that raises an interrupt, there must be a corresponding device interrupt handler.

For each device part of the hardware, the OS assigns a unique device identifier (devid) which is announced to the application programmer. (The device identifiers are specific to the particular installation). It is assumed that device identifiers are distinguishable from file handles. A user program can write a word into a device using the write system call. The read system call is used when the device allows a word to be read. Read and Write are the only system calls associated with devices.

The standard input and the standard output are two special stream devices with predefined identifiers STDIN = -1 and STDOUT = -2. Standard output permits only Write and standard input permits only Read. The Read operation typically puts the process executing the operation to sleep for console input from the user. When the user inputs data, the console device must send a hardware interrupt. The corresponding handler routine is called terminal handler. The terminal handler is responsible for waking up processes that are blocked for input console

Chapter 8

System Call Interface

Application programmers interact with the Operating System using the system calls. System calls are stored in the disk and are loaded into memory when the OS is loaded by the bootstrap loader. When a process invokes a system call, the process is interrupted and control goes to the corresponding interrupt service routine of the kernel, resulting in a switch from user mode to kernel mode. Once the system call is carried out, the control goes back to the application program, with a switch back to the user mode.

The system calls of eXpOS are classified into file system calls, process system calls and system calls for access control and synchronization.

8.1   File System Calls

(1) Create System Call

The Create operation takes as input a filename and creates an empty file by that name. If a root entry for the file already exists, then the system call returns 0 (success). Otherwise, it creates a root entry for the file name, sets the file type to DATA and file size to 0. Note that the file name must be a character string and must not be root.

(2) Delete System Call

Delete removes the file from the file system and removes its root entry. A

file that is currently opened by any application cannot be deleted. Root file also cannot be deleted.

(3) Open System Call

For a process to read/write a file, it must first open the file. Only data and root files can be opened. The Open operation returns a file descriptor. An application can open the same file several times and each time, a different descriptor will be returned by the Open operation. The file descriptor must be passed as argument to other file system calls, to identify the open instance of the file.

The OS associates a file pointer with every open instance of a file. The file pointer indicates the current location of file access (read/write). The Open system call sets the file pointer to 0 (beginning of the file).

(4) Close System Call

After all the operations are done, the user closes the file using the Close system call. The file descriptor ceases to be valid once the close system call is invoked.

(5) Read System Call

The file descriptor is used to identify an open instance of the file. The Read operation reads one word from the position pointed by the file pointer and stores it into the buffer. After each read operation, the file pointer advances to the next word in the file.

(6) Write System Call

The file descriptor is used to identify an open instance of the file. The Write operation writes the word stored in the buffer to the position pointed by

the file pointer of the file. After each Write operation, the file pointer advances to the next word in the file.

(7) Seek System Call

The Seek operation allows the application program to change the value of the file pointer so that subsequent Read/Write is performed from a new position in the file. The new value of the file pointer is determined by adding the offset to the current value. (A negative Offset will move the pointer backwards). An Offset of 0 will reset the pointer to the beginning of the file.

8.2 Process System Calls

(1) Fork System Call

Replicates the process invoking the system call. The heap, code and library regions of the parent are shared by the child. A new stack is allocated to the child and the parent's stack is copied into the child's stack.

When a process executes the Fork system call, the child process shares with the parent all the file and semaphore descriptors previously acquired by the parent. Semaphore/file descriptors acquired subsequent to the fork operation by either the child or the parent will be exclusive to the respective process and will not be shared.

(2) Exec System Call

Exec destroys the present process and loads the executable file given as input into a new memory address space. A successful Exec operation results in the extinction of the invoking application and hence never returns to it. All open instances of file and semaphores of the parent process are closed. However, the newly created process will inherit the PID of the

calling process.

(3) Exit System Call

Exit system call terminates the execution of the process which invoked it
and destroys its memory address space. The calling application ceases to
exist after the system call and hence the system call never returns.

(4) Getpid System Call

Returns the process identifier of the invoking process. The system call
does not fail.

(5) Getppid System Call

Returns to the calling process the value of the process identifier of its
parent. The system call does not fail.

(6) Shutdown System Call

Shutdown system call terminates all processes and halts the machine.

Arguments : None

Return Value : None

8.3   System calls for access control and synchronization

(1) Wait System Call

The current process is blocked till the process with PID given as argument
executes a Signal system call or exits. Note that the system call will fail
if a process attempts to wait for itself.

(2) Signal System Call

All processes waiting for the signaling process are resumed. The system
call does not fail.

(3) FLock System Call

To lock a file so that other applications running concurrently are not permitted to access the file till the calling process unlocks it. If the file is already locked by some other process, the system call waits for the file to be unlocked, locks it, and returns to the calling process.

(4) FUnLock System Call

FUnLock operation allows an application program to unlock a file which the application had locked earlier, so that other applications are no longer restricted from accessing the file.

(5) Semget System Call

This system call is used to obtain a binary semaphore. eXpOS has a fixed number of semaphores. The calling process can share the semaphore with its child processes using the fork system call.

(6) Semrelease System Call

This system call is used to release a semaphore descriptor held by the process.

(7) SemLock System Call

This system call is used to lock the semaphore. If the semaphore is already locked by some other process, then the calling process goes to sleep and wakes up only when the semaphore is unlocked. Otherwise, it locks the semaphore and continues execution.

(8) SemUnLock System Call

This system call is used to unlock a semaphore that was previously locked by the calling process. It wakes up all the processes which went to sleep trying to lock the semaphore while the semaphore was locked by the calling

process.

Chapter 9

Design of Data Structures

Data Structures can be classified into - Memory Data Structures (In-core) and Disk Data Structures. The Disk Data Structures are loaded to memory by the OS startup code and stored back when system terminates.

9.1 Disk Data Structures

9.1.1 Inode Table

The Inode table is stored in the disk and has an entry for each file present in the disk (A copy of the Inode table is maintained in the memory when the OS is running). It consists of 32 entries. Thus eXpFS permits a maximum of 32 files. The entry of an Inode table has the following format:

FILE TYPE	FILE NAME	FILE SIZE	DATA BLOCK 1	DATA BLOCK 2	DATA BLOCK 3	DATA BLOCK 4	Unused

Figure 9.1: Inode Table

- FILE TYPE (1 word) - specifies the type of the given file (ROOT indicated by 1 , DATA indicated by 2 or EXEC indicated by 3). More information about file types is given here.

- FILE NAME (1 word) - Name of the file

- FILE SIZE (1 word) - Size of the file. Maximum size for File = 2048 words

- DATA BLOCK 1 to 4 (4 words) - each DATA BLOCK column stores the block number of a data block of the file. If a file does not use a particular DATA BLOCK , it is set to -1.

- Unused (9 words)

All unused entries are set to -1.

9.1.2   Disk Free List

The Disk Free List consists of 512 entries. Each entry is of size one word. For each block in the disk there is an entry in the Disk Free List which contains either 0 (free) or a number indicating the number of processes sharing the block.

9.1.3   Root File

The Root File is stored in the disk and has an entry for each file present in the disk (A copy of the Root File is maintained in the memory when the OS is running). It consists of 32 entries.

The entry of Root file has the following format:

FILE NAME	FILE SIZE	FILE TYPE	Unused

Figure 9.2: Root File

- FILE NAME (1 word) - Name of the file

- FILE SIZE (1 word) - Size of the file

- FILE TYPE (1 word) - specifies the type of the given file (ROOT indicated by 1 , DATA indicated by 2 or EXEC indicated by 3).

- Unused (5 words)

All unused entries are set to -1.

9.2   Memory Data Structures

9.2.1   Process Table

The Process Table (PT) contains an entry for each process present in the system. The entry is created when the process is created by a Fork system call. Each entry contains several fields that stores all the information pertaining to a single process. The maximum number of entries in PT (which is maximum number of processes allowed to exist at a single point of time in eXpOS) is 32.

Any entry of PT has the following fields:

TICK	PID	PPID	STATE	PER-PROCESS RESOURCE TABLE	INODE INDEX	INPUT BUFFER	MODE FLAG	KERNEL RE-ENTRY POINT	PTBR	MACHINE STATE POINTER	Unused

Figure 9.3: Process Table

- TICK (1 word)- keeps track of how long the process was in memory. It has an initial value of 0 and is updated whenever the scheduler is called. TICK is reset to 0 when a process is swapped out or in.

- PID (1 word) - process descriptor, a number that is unique to each process. This field is set by Fork System Call.

- PPID (1 word) - process descriptor of the parent process. This field is set by Fork System Call.

- STATE (2 words) - a two tuple that describes the current state of the process. PER-PROCESS RESOURCE TABLE (16 words) - Per-Process Resource Table contains information about the files opened by the process as well as semaphores used by the process.

- INODE INDEX (1 word)- Pointer to the Inode entry of the executable file, which was loaded into the process's address space.

- INPUT BUFFER (1 word) - Buffer used to store the input read from the terminal. Whenever a word is read from the terminal, Terminal Interrupt Handler will store the word into this buffer.

- MODE FLAG (1 word) - Used to indicate whether the process is executing in kernel or user mode. It is set to 1 when it enters the kernel mode (upon entering a system call) and resets to 0 when it returns back to the user mode (upon returning from a system call).

- KERNEL RE-ENTRY POINT (1 word) - Contains the return address when a process voluntarily schedules out itself inside a blocking system call.

- PTBR (1 word) - pointer to PER PROCESS PAGE TABLE.

- MACHINE STATE POINTER(1 word) - pointer to a structure that gives the machine state when the process was last executed. The scheduler uses this structure to store the context of the process while scheduling. MACHINE STATE is architecture-dependent.

- Unused (5 words)

Invalid entries are represented by -1.

States

The tuple can take the following values

- (RUNNING,_): The process is in execution. This field is set by Scheduler when a process is scheduled.

- (READY,_): The process is ready to be scheduled. (WAIT_PROCESS,WAIT_PID): The process is waiting for a signal from another process whose

- PID is WAIT_PID. This field is set by WAIT system call.

- (WAIT_FILE, FTENTRY): The process is blocked for a file whose file table entry index is FTENTRY.

- (WAIT_DISK,_): The process is blocked because of one of the following reasons: a) It is waiting for disk to complete a disk-memory transfer operation it had initiated OR b) It wants to execute a disk transfer, but the disk is busy, handling a disk-memory transfer request issued by some other process.

- (WAIT_SEMAPHORE,SEMID): The process is waiting for a semaphore (whose descriptor is SEMID) that was locked by some other process.

- (WAIT_MEM, _): The process is blocked due to unavailability of memory pages.

- (SWAPPED,_): The stack page of the process has been swapped out into disk.

- (SWAPPED_WAIT, WAIT_PID): The stack page of a process which was in (WAIT_PROCESS,WAIT_PID) state has been swapped out into disk

- (WAIT_BUFFER, BUFFERID): The process is waiting for disk buffer of index BUFFERID to be unlocked.

- (WAIT_TERMINAL, _): The process is waiting for Read from Terminal to be completed.

Machine State

The OS scheduler can schedule out a process in execution and save its machine state here. Later, the scheduler can restore the machine state to start ex-

ecution from where it had stopped. The initial values of this structure are fixed by the OS loader (or exec system call) and the values get updated whenever the process is preempted.

It would contain the following information:

SP	BP	IP	PTBR	PTLR	REGISTERS	Unused

Figure 9.4: Machine State

- SP (1 word)- Contents of Stack Pointer Register.

- BP (1 word) - Contents of Base Pointer Register.

- IP (1 word) - Contents of Instruction Pointer Register.

- PTBR (1 word)- Contents of Page Table Base Register.

- PTLR (1 word)- Contents of Page Table Length Register.

- REGISTERS (24 words) - Machine registers.

- Unused (3 words)

Per Process Resource Table

The Per-Process Resource Table has 8 entries and each entry is of 2 words. For every instance of file opened (or a semaphore acquired) by the process, it stores the index of the File Table (or Semaphore Table) entry for the file (or semaphore). In case of a file, the second word stores the LSEEK position of the open instance of that file. The LSEEK position indicates the location in the file where a word is read from or written to corresponding to the open instance of the file. In case of a semaphore, the second word is unused.

The per-process resource table entry has the following format:

Index of File Table/ Semaphore Table entry (1 word)	LSEEK (1 word)

Figure 9.5: Per Process Resource Table

File descriptor, returned by Open system call, is the index of the per-process resource table entry for that open instance of the file. **Per Process Page Table**

The Per-Process Page Table contains information regarding the physical location of the pages of a process. Each valid entry of a page table stores the physical page number corresponding to each logical (virtual) page associated with the process. The logical page number can vary from 0 to 7 for each process. Therefore, each process has 8 entries in the page table.

Associated with each page table entry, typically auxiliary information is also stored. This is to store information like whether the process has write permission to the page, whether the page is dirty, referenced, valid etc. The exact details are architecture dependent.

PHYSICAL PAGE NUMBER	REFERENCE BIT	VALID BIT	WRITE BIT

Figure 9.6: Per Process Page Table

- Reference bit - The reference bit for a page table entry is set to 0 by the OS when the page is loaded to memory and the page table initialized. When a page is accessed by a running process, the corresponding reference bit is set to 1 by the machine. This bit is used by the page replacement algorithm of the OS.

- Valid bit - This bit is set to 1 by the OS when the physical page number

field of a page table entry is valid (i.e, the page is loaded in memory). It is set to 0 if the entry is invalid. The OS expects the architecture to generate a page fault if any process attempts to access an invalid page.

- Write bit - This bit is set to 1 by the OS if the page can be written and is set to 0 otherwise. The OS expects the architecture to generate an exception if any process attempts to modify a page whose write bit is not set.

9.2.2   File Table

The File Table stores the information about all the files that are open while the OS is running. It consists of 32 entries. Thus, there can be at most 32 open files in the system at any time.

Every entry of the File Table has the following format :

INODE INDEX	FILE OPEN COUNT	ULOCK	KLOCK	Unused

Figure 9.7: File Table

- INODE INDEX (1 word) - specifies the index of the entry for the file in the Inode table.

- FILE OPEN COUNT (1 word) - specifies the number of open instances of the file

- ULOCK (1 word)- If the file was locked by a process using FLock system call, this field specifies the PID of the process, otherwise, it is set to -1.

- KLOCK (1 word)- If the file is locked by the process inside a system call, this field specifies the PID of the process, otherwise, it is set to -1.

- Unused (4 words)

All invalid entries are set to -1.

9.2.3   Semaphore Table

The Semaphore Table contains details about all the semaphores used by the processes. It consists of 32 entries. Thus, there can be at most 32 semaphores used in the system at any time. For every semaphore entry in the per-process resource table, there is a corresponding entry in the semaphore table.

LOCKING PID	PROCESS COUNT	Unused

Figure 9.8: Semaphore Table

- LOCKING PID (1 word)- specifies the PID of the process which has locked the semaphore

- PROCESS COUNT (1 word) - specifies the number of processes which are sharing the semaphore.

- Unused (2 words)

All invalid entries are set to -1.

9.2.4   Buffer Table

The buffer table is a data structure that stores the information regarding the disk block stored in each buffer. The present version of eXpOS sets MAX_BUFFER = 4. Each buffer is identified by its index in the buffer table.
Each entry of the Buffer Table has the following format:

BLOCK NUMBER	DIRTY BIT	LOCKING PID	Unused

Figure 9.9: Buffer Table

- BLOCK NUMBER (1 word) - specifies block number of the disk block which is currently stored in the buffer. It is set to -1 if the buffer does not contain any valid disk block.

- DIRTY BIT (1 word) - specifies whether the block stored in the buffer has been modified. It is set to 0 when a disk block is loaded into the buffer and is set to 1 when modified by the Write System Call.

- LOCKING PID (1 word) - specifies the PID of the process which has currently locked the buffer while executing a file system call.

- Unused - (1 word)

Free entries are represented by -1 in all the fields.

9.2.5 Disk Status Table

The Disk Status Table keeps track of these disk-memory transfers. Every time a disk operation is invoked, the information regarding the operation like the disk block and the memory page involved, the process that invoked the operation and type of disk operation are stored in Disk Status Table by the OS.

STATUS	LOAD/STORE BIT	PAGE NUMBER	BLOCK NUMBER	PID	Unused

Figure 9.10: Disk Status Table

- STATUS (1 word) - specifies whether the disk is free (indicated by 0) or busy (indicated by 1) handling a memory-disk transfer.

- LOAD/STORE BIT (1 word) - specifies whether the operation being done on the device is a load (indicated by 0) or store (indicated by 1).

- PAGE NUMBER (1 word) - specifies the memory page number involved in the disk transfer

- BLOCK NUMBER (1 word) - specifies the disk block number involved in the disk transfer.

- PID (1 word) - specifies the PID of the process which invoked the disk transfer. It is set to -PID if the disk transfer was initiated by the scheduler, to load/store a page of the process with identifier PID.

- Unused (3 words)

9.2.6   System Status Table

It keeps the information about the number of free pages in memory, the number of processes blocked because memory is unavailable and, the number of processes in swapped state.

The System Status Table has the following format:

LOCKING PID	PROCESS COUNT	Unused

Figure 9.11: System Status Table

- MEM_FREE_COUNT (1 word) - specifies the number of free pages available in memory.

- WAIT_MEM_COUNT (1 word) - specifies the number of processes waiting for memory.

- SWAPPED_COUNT (1 word)- specifies the number of processes in (SWAPPED, _) or (SWAPPED_WAIT, _) state.

- Unused (1 word)

9.2.7 Terminal Status Table

The Terminal Status Table keeps track of the Read operations done on the terminal. Every time a Read system call is invoked on the terminal, the PID of the process that invoked the system call is stored in Terminal Status Table. Every entry of the Terminal Status Table has the following format:

MEM_FREE_COUNT	WAIT_MEM_COUNT	SWAPPED_COUNT	Unused

Figure 9.12: Terminal Status Table

- STATUS (1 word) - specifies whether the terminal is free or is being used by a process to read input. This field is initially set to 0. It is changed to 1 whenever terminal is busy. The Terminal Interrupt Handler sets back the status to 0 upon completion of Terminal Read.

- PID (1 word) - specifies the PID of the process which is currently reading input from the terminal. This field is invalid when STATUS is 0.

- Unused (2 words)

9.2.8 Memory Free List

The Memory free list is a data structure used for keeping track of used and unused pages in the memory. It consists of 128 entries and each entry is of size one word.

Each entry of the free list contains either the value 0 indicating that the

corresponding page in the memory is free (not allocated to any process) or contains the number of processes that are sharing the page.

Chapter 10

Algorithms Used

Algorithm 1 Create System Call
___
Input: Filename
Output: 0 (Success) or -1 (No Space for file)
    If the file is present in the system, return 0.
    Find the index of a free entry in the Inode Table. If no free entry found, return -1.
    Allocate the Inode Table entry to the file.
    Update the Root file by adding an entry for the new file.
    return 0

Algorithm 2 Delete System Call
___
Input: Filename
Output: 0 (Success) or -1 (File not found) or -2 (File is open)
    If file is not present in Inode Table or if it is not a DATA file, return -1.
    Find the index of the file in the Inode Table.
    If File Table entry exists with the same index as found above, return -2.
    Update Inode Table.
    Update the Root file by invalidating the root entry for the file.
    return 0

Algorithm 3 Open System Call

Input: Filename

Output: File Descriptor (Success) or -1 (File not found) or -2 (Process has reached
  its limit of resources) or -3 (System has reached its limit of open files)

  If file is not present in Inode Table or if it is of type EXEC, return -1.

  Find the index of the Inode Table entry of the file.

  Allocate entry in Per Process Resource Table

  if file is already open **then**

      Get the File Table entry of the file and increment the File Open Count.

  else

      Find a free entry in the File Table. If there are no free entries, return -3.

  end if

  Allocate the Per-Process Resource Table entry to the file.

  Update the File Table entry.

  return from system call with the index of the Per-Process Resource Table entry.

Algorithm 4 Close System Call

Input: File Descriptor

Output: 0 (Success) or -1 (File Descriptor is invalid) or -2 (File is locked by calling
  process)

  If file descriptor does not correspond to valid entry in Per Process Resource
  Table, return -1

  Get the index of the File Table entry from Per-Process Resource Table entry.

  If file is locked by the current process, return -2.

  In the File Table Entry found above, decrement the File Open Count. If the
  count becomes 0, invalidate the entry.

  Invalidate the Per-Process Resource Table entry.

  return 0

Algorithm 5 Read System Call

Input: File Descriptor and a Buffer (a String/Integer variable) into which a word is to be read from the file

Output: 0 (Success) or -1 (File Descriptor is invalid) or -2 (File pointer has reached the end of file)

    if input is to be read from terminal **then**

        Wait till terminal gets the input for the current process.

        Copy the word from the input buffer of the Process Table to the buffer passed as argument and return with 0.

    end if

    If file descriptor does not correspond to a valid entry in Per Process Resource Table, return -1.

    Check the validity of the File pointer

    Find the block and the position in the block from which input is read.

    while the file is locked by a process other than the current process **do**

        put the process to sleep and call the scheduler.

    end while

    Lock the file.

    Get the buffer page number from block number.

    while the buffer is locked by a process other than the current process **do**

        put the process to sleep and call the scheduler.

    end while

    Lock the buffer page.

    if the buffer contains a block other than the required block **then**

        Bring the block to the buffer from the disk

    end if

    Copy the word at the offset position of the block into the buffer passed as argument and Increment lseek position.

    Unlock the buffer and wake up all processes waiting for the buffer.

    Unlock the file and wake up all processes waiting for the file.

    return 0

Algorithm 6 Write System Call

Input: File Descriptor and a Word to be written

Output: 0 (Success) or -1 (File Descriptor given is invalid) or -2 (No disk space)

  if word is to be written to standard output **then**

    Issue the machine instruction to output the given word and return 0.

  end if

  If file descriptor does not correspond to valid entry in Per Process Resource Table, return -1.

  Get the index of the File Table entry and lseek position from the Per Process Resource Table entry.

  If lseek is equal to MAX_FILE_SIZE - 1, return -2.

  Find the block and position in the block to which data has to be written.

  while the file is locked by a process other than the current process **do**

    put the process to sleep and call the scheduler.

  end while

  Lock the file.

  Get the buffer page number from block number.

  while the buffer is locked by a process other than the current process **do**

    put the process to sleep and call the scheduler.

  end while

  Lock the buffer page.

  if the buffer contains a block other than the required block **then**

    Bring the block to the buffer from the disk

  end if

  Copy the word passed as argument to the offset position in the buffer.

  Set dirty bit to 1 in the Buffer Table entry and increment lseek position.

  Increment file size in the Inode Table entry and Root file entry.

  Unlock the buffer and wake up all processes waiting for buffer.

  Unlock the file and wake up all processes waiting for the file.

  return 0.

Algorithm 7 Seek System Call

Input: File Descriptor and Offset

Output: 0 (Success) or -1 (File Descriptor given is invalid) or -2 (Offset value moves the file pointer to a position outside the file)

If file descriptor does not correspond to valid entry in Per Process Resource Table, return -1.

Get the index of the File Table entry and lseek position from the Per Process Resource Table entry.

Check the validity of the given offset

if given offset is 0 **then**

    Set lseek value in the Per-Process Resource Table entry to 0.

else

    Change the lseek value in the Per-Process Resource Table entry to lseek+offset.

end if

return 0

Algorithm 8 Fork System Call

Input: None

Output: Process Identifier to the parent process and 0 to child process (Success) or -1 (Number of processes has reached its maximum, returned to parent)

If no free entry in the Process Table, return -1.

Find the index of a free entry in the Process Table and set the PID and PPID field.

Count the number of valid stack pages from the Page Table of the parent process.

If sufficient number of free pages are not present in memory, then increment the WAIT_MEM_COUNT in the System Status Table.

while sufficient number of free pages are not present in memory **do**

    put the process to sleep and call the scheduler.

end while

for each stack page of the parent process **do**

    if the stack page is valid **then**

        Allocate a free page to the child.

        Copy the stack page of the parent into the child stack page.

    else

        Share the swap block containing the stack page with the child.

    end if

end for

Construct the context of the child process.

Set the return value to 0 for the child process.

The PID of the child process is set as the return value for the parent process.

return to the parent process.

Algorithm 9 Exec System Call

Input: Filename

Output: -1 (File not found or file is of invalid type)

 If file not found in system or file type is not EXEC, return -1.

 For each page of the current process that is swapped out, find the swap block and decrement its entry in the Disk Free List.

 Setup code and library pages.

 Free the heap pages

 In the Process Table entry of the current process, set the Inode Index field to the index of Inode Table entry for the file.

 Close all files opened by the current process.

 Release all semaphores held by the current process.

 Set SP to the start of the stack region and IP to the start of the code region.

 return from system call to newly loaded process.

Algorithm 10 Exit System Call

Input: None

Output: None

 If no more processes to schedule, shutdown the machine.

 Unlock and close all files opened by the current process.

 Release all the semaphores used by the current process.

 Wake up all processes waiting for the current process.

 Invalidate the Process Table entry and the page table entries of the current process

 Invoke the scheduler to schedule the next process.

Algorithm 11 Getpid System Call

Input: None

Output: Process Identifier (Success)

 Find the PID of the current process from the Process Table.

 return the PID of current process.

Algorithm 12 Getppid System Call

Input: None

Output: Process Identifier (Success)

 Find the PPID of the current process from the Process Table.

 return PPID.

Algorithm 13 Wait System Call

Input: Process Identifier of the process for which the current process has to wait.

Output: 0 (Success) or -1 (Given process identifier is invalid or it is the pid of the invoking process)

If process is intending to wait for itself or for a non-existent process, return -1.

Change the status from (RUNNING, _) to (WAIT_PROCESS, Argument_PID) in the Process Table.

Invoke the Scheduler to schedule the next process.

return 0

Algorithm 14 Signal System Call

Input: None

Output: 0 (Success)

Wake up all processes waiting for the current process.

return 0

Algorithm 15 FLock System Call

Input: File Descriptor

Output: 0 (Success) or -1 (File Descriptor is invalid)

If file descriptor does not correspond to valid entry in Per Process Resource Table, return -1.

while the file is locked by a process other than the current process **do**

Change the state to (WAIT_FILE, ftindex) where ftindex is the File Table index of the locked file and call the Scheduler.

end while

Change the ULock field of the file table to PID of current process.

return 0

Algorithm 16 FUnLock System Call

Input: File Descriptor

Output: 0 (Success) or -1 (File Descriptor is invalid) or -2 (File was not locked by the calling process)

If file descriptor does not correspond to a valid entry in Per Process Resource Table, return -1.

if file is locked **then**

If current process has not locked the file, return -2.

Unlock the file.

Wake up all processes waiting for the file.

end if

return 0

Algorithm 17 Semget System Call

Input: None

Output: semaphore descriptor (Success) or -1 (Process has reached its limit of resources) or -2 (Number of semaphores has reached its maximum)

Find the index of a free entry in Per Process Resource Table. If no free entry, then return -1.

Find the index of a free entry in Semaphore table and increment the process count. If no free entry, return -2.

Store the index of the Semaphore table entry in the Per Process Resource Table entry.

return Semaphore Table entry index.

Algorithm 18 Semrelease System Call

Input: Semaphore Descriptor

Output: 0 (Success) or -1 (Semaphore Descriptor is invalid)

If Semaphore descriptor is not valid or the entry in Per Process Resource Table is not valid , return -1.

if semaphore is locked by the current process **then**

    Unlock the semaphore before release.

end if

Decrement the process count of the semaphore.

Invalidate the Per-Process resource table entry.

return 0

Algorithm 19 SemLock System Call

Input: Semaphore Descriptor

Output: 0 (Success or the semaphore is already locked by the current process) or -1 (Semaphore Descriptor is invalid)

If Semaphore descriptor is not valid or the entry in Per Process Resource Table is not valid , return -1.

while the semaphore is locked by a process other than the current process **do**

    Put the current process to sleep

    Call scheduler

end while

Change the Locking PID to PID of the current process in the Semaphore Table .

return 0

Algorithm 20 SemUnLock System Call

Input: Semaphore Descriptor

Output: 0 (Success) or -1 (Semaphore Descriptor is invalid) or -2 (Semaphore was not locked by the calling process)

    If Semaphore descriptor not valid or the entry in Per Process Resource Table is not valid , return -1.

    if semaphore is locked **then**

        If current process has not locked the semaphore, return -2.

        Unlock the semaphore.

        Wake up all processes waiting for the semaphore.

    end if

    return 0

Algorithm 21 Timer Interrupt Handler

    Save the context of current process and mark it as ready.

    Increment TICK

    if free memory pages present **then**

        if there are sleeping processes that requires memory pages **then**

            Wake up all processes that requires memory pages.

        else

            if disk is free and there are swapped processes **then**

                Swap in the seniormost swapped out process.

            end if

        end if

    else

        if disk is free and there are processes that require memory pages **then**

            Use the modified second chance algorithm to find an unreferenced page

            if the unreferenced page is a code page **then**

                In the page table entry of the unreferenced page, store the code block number.

            else

                If the unreferenced page is a stack or heap page, swap the page to a free swap block in the disk.

            end if

            if unreferenced page is pointed to by the stack pointer of the process **then**

                Mark the process that owns the page as swapped out.

            end if

        end if

    end if

    Schedule the next ready process

Algorithm 22 Disk Interrupt Handler

In the Disk Status Table, set the STATUS field to 0, indicating that the disk is no longer busy.

if load/store was issued by a process **then**

    Wake up all processes waiting for the disk.

    If the loaded block was a swap block, decrement the corresponding Disk Free List entry.

else

    Get the PID of the process for which the scheduler had issued the load/store instruction.

    if the disk operation was a load operation **then**

        Update Disk Free List.

        Update the Page Table of the process.

        Mark the process as ready.

        Decrement SWAPPED_COUNT in System Status Table.

    else

        Update Memory Free List.

        Update Page Table of the process.

        Increment the MEM_FREE_COUNT in the System Status Table.

    end if

end if

Algorithm 23 Exception Handler

If the exception is not caused by a page fault, display the cause and exit the process.

Find the page table entry of the page causing the exception.

If there are no free pages in the memory, then increment the WAIT_MEM_COUNT in the System Status Table.

while there are no free pages in memory **do**

    Put the current process to sleep and call scheduler

end while

Find a free page in memory.

if the page that caused exception is stored in the disk **then**

    while the disk is busy **do**

        Put the current process to sleep and call scheduler

    end while

    Load the block to the memory page found above.

    Put the current process to sleep and call scheduler.

end if

Update the Page Table entry of the page that caused exception.

return to the user process.

Algorithm 24 Terminal Interrupt Handler

Set the status field in Terminal Status Table to 0.
Locate the Process Table entry of the process that read the data.
Copy the word read from standard input to the Input Buffer field in the Process
Table entry.
Set the PID field of Terminal Status Table to -1.
Wake up all processes waiting for terminal.

Algorithm 25 OS Startup Code

Load the software interrupts and handler routines to memory.
Load the init program to memory and set up its machine context.
Load the idle program to memory and set up its machine context.
Load the disk data structures to the memory.
Initialize all the memory data structures.
Schedule the init process for execution.

Algorithm 26 Shutdown System Call

while disk is not free **do**
    Put the current process to sleep and call scheduler
end while
Store Inode Table to the disk.
Store dirty pages contained in buffer to the disk.
Store Disk Free List to the disk.
Halt the machine.

Algorithm 27 Shell process

while true **do**
    Get the program to be executed using read system call
    If the shutdown instruction is received, invoke shutdown system call.
    childPID = fork();
    if childPID == 0 **then**
        Execute the program given by read.
    else
        Wait for child to finish execution.
    end if
end while

Algorithm 28 Idle process

while true **do**
end while

Chapter 11

Work Done

∗ The existing OS data structures were redesigned to incorporate the changes done.

∗ New data structures like Buffer Table, Disk Status Table, Semaphore Table, System Status Table etc were designed.

∗ System calls were redesigned to incorporate asynchronous operations, buffer cache and pre-emptive scheduling.

∗ Directory structure was introduced in eXpFS.

∗ Executable file format was designed.

∗ Algorithms for system calls and other interrupt handlers were designed.

∗ Webpage for Project eXpOS was created. *http : //exposnitc.github.io/*

Chapter 12

Future Work

* Building a Roadmap for guidance.

* Implementing and testing the algorithms designed

* Adding more features like file permissions and Super-User previleges.

* Adding Multi-Threading feature.

Chapter 13

Conclusion

This project aims to create a simpler version of an operating system which allows students to acquire insight into the working of a real operating system.

Bibliography

Appendix A

Appendix

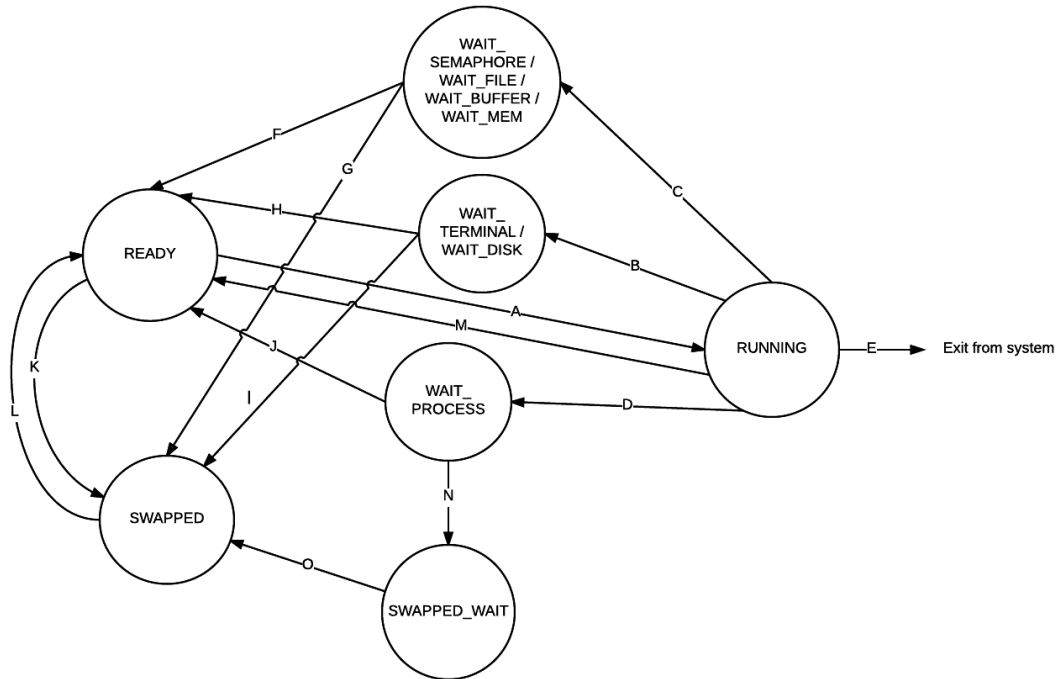## A.1   State Transition Diagram in eXpOS

Figure A.1: State Transitions

A

READY → RUNNING : The scheduler has scheduled the process for execution.

B

RUNNING → WAIT_TERMINAL : The process is either waiting for access to the terminal or for data to be inputted through terminal.

RUNNING → WAIT_DISK : The process is either waiting for access to the disk or for the disk operation to finish.

C

RUNNING → WAIT_SEMAPHORE : The semaphore which the process is trying to use, is found to be locked.

RUNNING → WAIT_FILE : The file which the process is trying to read/write, is found to be locked.

RUNNING → WAIT_BUFFER : The buffer which the process is trying to use, is found to be locked.

RUNNING → WAIT_MEM : The process requires a free memory page but there are none in the memory.

D

RUNNING → WAIT_PROCESS : The process is waiting for another process to either exit or to signal it.

E

RUNNING → Exit from system : The process has either completed execution or has invoked an Exit System Call.

F

WAIT_SEMAPHORE → READY : The semaphore for which the process

was waiting, is now unlocked.

WAIT_FILE → READY : The file for which the process was waiting, is now unlocked.

WAIT_BUFFER → READY : The buffer for which the process was waiting, is now unlocked.

WAIT_MEM → READY : There are free pages in memory.

G

WAIT_SEMAPHORE / WAIT_FILE / WAIT_BUFFER / WAIT_MEM → SWAPPED : The current stack page of the process has been swapped out.

H

WAIT_TERMINAL → READY : The input data has been read from terminal and terminal is free to be used by any process.

WAIT_DISK → READY : The disk operation is complete.

I

WAIT_TERMINAL / WAIT_DISK → SWAPPED : The current stack page of the process has been swapped out.

J

WAIT_PROCESS → READY : The process has either received a signal from the process it was waiting for or the latter has exited the system.

K

READY → SWAPPED : The current stack page of the process has been swapped out.

L

SWAPPED → READY : The stack page required by the process to continue execution was swapped in.

M

RUNNING → READY : Context switch caused by the timer interrupt routine.

N

WAIT_PROCESS → SWAPPED_WAIT : The current stack page of the process was swapped out while waiting for another process.

O

SWAPPED_WAIT → SWAPPED : The process in SWAPPED_WAIT state has either received a signal from the process it was waiting for or the latter has exited.